



ua

Universidade de Aveiro

Mestrado em Engenharia de Computadores e Telemática

Arquitecturas de Alto Desempenho

DLX – Pipelining 2

Academic year 2021/2022 Adaptation of exercise guide by Nuno Lau/José Luís Azevedo

4. Write a program that orders the values of an integer array stored in memory from the largest value to the smallest one.

Original code

```
.data
values: .word 1,2,3,4,5,6,7,8,9,10      ; values to be ordered
nelem:  .word 10                        ; array size
.text
.global main
main:   addi r1,r0,nelem                ; r1 = add(nelem)
        lw   r1,0(r1)                  ; r1 = val(nelem)
        addi r2,r0,values              ; r2 = add(values[0])
        add  r3,r0,r0                  ; r3 = i = 0 (counting variable)
        addi r8,r1,-1                  ; r8 = nelem - 1
loop1:  slt  r9,r3,r8                  ; r9 = (i < nelem - 1)
        beqz r9,end                   ; is the end of operations been
                                           ; reached?
        addi r6,r2,4                   ; r6 = add(values[j])
        lw   r4,0(r2)                  ; r4 = val(values[i])
        addi r5,r3,1                   ; r5 = j = i+1 (counting variable)
loop2:  lw   r7,0(r6)                  ; r7 = val(values[j])
        slt  r9,r4,r7                  ; r9 =
                                           ; (val(values[i]) < val(values[j]))
        beqz r9,goon                  ; no element swap is required?
        add  r9,r4,r0                  ; r9 = tmp = val(values[i])
        add  r4,r7,r0                  ; val(values[i]) = val(values[j])
        add  r7,r9,r0                  ; val(values[j]) = tmp
        sw   0(r2),r4                  ; val(values[i]) = r4
        sw   0(r6),r7                  ; val(values[j]) = r7
goon:   addi r5,r5,1                   ; j = j + 1
        addi r6,r6,4                   ; r6 = add(values[j])
        slt  r9,r5,r1                  ; r9 = (j < nelem)
        bnez r9,loop2                 ; are still elements to be
                                           ; compared?
        addi r3,r3,1                   ; i = i + 1
        addi r2,r2,4                   ; r2 = add(values[i])
        j    loop1                     ; check for end of operations
end:    trap 0                         ; end of program
```

- 4.1. Run the code in the DLX simulator with the *forwarding* option turned off and, through the interspersing of *nop* instructions to take care of the hazards, make it execute correctly. What is the number of clock cycles for the full execution? What is the speed up attained relatively to the original code being run in a non-pipelined processor with a clock cycle time five times longer? Take into consideration that in the latter case not all instructions take the same time to execute.

For comparing purposes, one assumes that the non-pipelined processor executes *store* and *branch/jump* instructions in 4 equivalent clock cycles and all other instructions in 5 equivalent clock cycles of the pipelined processor.

Code with `nop` instructions interspersed to make it run as intended

```

        .data
values:  .word  1,2,3,4,5,6,7,8,9,10      ; values to be ordered
nelem:   .word  10                        ; array size
        .text
main:    .global  main
        addi    r1,r0,nelem      ; r1 = add(nelem)
        nop
        nop
        lw      r1,0(r1)         ; r1 = val(nelem)
        addi    r2,r0,values     ; r2 = add(values[0])
        add     r3,r0,r0         ; r3 = i = 0 (counting variable)
        addi    r8,r1,-1         ; r8 = nelem - 1
        nop
        nop
loop1:   slt     r9,r3,r8         ; r9 = (i < nelem - 1)
        nop
        nop
        beqz    r9,end          ; is the end of operations been
                                ; reached?
        nop
        nop
        nop
        addi    r6,r2,4          ; r6 = add(values[j])
        lw      r4,0(r2)         ; r4 = val(values[i])
        addi    r5,r3,1          ; r5 = j = i+1 (counting variable)
loop2:   lw      r7,0(r6)         ; r7 = val(values[j])
        nop
        nop
        slt     r9,r4,r7         ; r9 =
                                ; (val(values[i]) < val(values[j]))
        nop
        nop
        beqz    r9,goon         ; no element swap is required?
        nop
        nop
        nop
        add     r9,r4,r0         ; r9 = tmp = val(values[i])
        add     r4,r7,r0         ; val(values[i]) = val(values[j])
        add     r7,r9,r0         ; val(values[j]) = tmp
        nop
        sw      0(r2),r4         ; val(values[i]) = r4
        nop
        sw      0(r6),r7         ; val(values[j]) = r7
goon:    addi    r5,r5,1          ; j = j + 1
        nop
        addi    r6,r6,4          ; r6 = add(values[j])
        slt     r9,r5,r1         ; r9 = (j < nelem)
        nop
        nop
        bnez    r9,loop2        ; are still elements to be
                                ; compared?
        nop
        nop
        nop
        addi    r3,r3,1          ; i = i + 1
        addi    r2,r2,4          ; r2 = add(values[i])
        j       loop1           ; check for end of operations
        nop
        nop
        nop
end:     trap    0               ; end of program

```

Putting aside the `nop` instructions, 620 instructions are executed in 1380 clock cycles

$$\begin{aligned}\text{instruction count} &= 5 + 8 \cdot 9 + 12 \cdot \sum_{i=1}^9 i + 3 = \\ &= 5 + 72 + 540 + 3 = 620\end{aligned}$$

$$\text{instruction count}_{store} = 2 \cdot \sum_{i=1}^9 i = 90$$

$$\begin{aligned}\text{instruction count}_{br / jmp} &= 2 \cdot 9 + 2 \cdot \sum_{i=1}^9 i + 1 = \\ &= 18 + 90 + 1 = 109.\end{aligned}$$

Thus, the equivalent number of clock cycles for the non-pipelined processor is

$$\begin{aligned}\text{clock cycles}_{nonpipe} &= (\text{instruction count}_{store} + \text{instruction count}_{br / jmp}) \cdot 4 + \\ &\quad + \text{instruction count}_{other} \cdot 5 = \\ &= 199 \cdot 4 + 421 \cdot 5 = 2901\end{aligned}$$

and the speed up is given by

$$\text{speed up} = \frac{\text{clock cycles}_{nonpipe}}{\text{clock cycles}_{pipe}} = \frac{2901}{1380} = 2,10.$$

4.2. Turn the *forwarding* option on and discard the `nop` instructions that are no longer required.
What was the speed up now attained?

With the *forwarding* option turned on, most of the `nop` instructions can be removed.

```

.data
values: .word 1,2,3,4,5,6,7,8,9,10      ; values to be ordered
nelem:  .word 10                        ; array size
.text
.global main
main:   addi r1,r0,nelem                ; r1 = add(nelem)
        lw   r1,0(r1)                  ; r1 = val(nelem)
        addi r2,r0,values               ; r2 = add(values[0])
        add  r3,r0,r0                   ; r3 = i = 0 (counting variable)
        addi r8,r1,-1                   ; r8 = nelem - 1
loop1:  slt  r9,r3,r8                   ; r9 = (i < nelem - 1)
        beqz r9,end                    ; is the end of operations been
                                           ; reached?

        nop
        nop
        addi r6,r2,4                    ; r6 = add(values[j])
        lw   r4,0(r2)                   ; r4 = val(values[i])
        addi r5,r3,1                    ; r5 = j = i+1 (counting variable)
loop2:  lw   r7,0(r6)                   ; r7 = val(values[j])
        nop
        slt  r9,r4,r7                   ; r9 =
                                           ; (val(values[i]) < val(values[j]))
        beqz r9,goon                    ; no element swap is required?

        nop
        nop
        add  r9,r4,r0                   ; r9 = tmp = val(values[i])
        add  r4,r7,r0                   ; val(values[i]) = val(values[j])
        add  r7,r9,r0                   ; val(values[j]) = tmp
        sw   0(r2),r4                   ; val(values[i]) = r4
        sw   0(r6),r7                   ; val(values[j]) = r7
goon:   addi r5,r5,1                     ; j = j + 1
        addi r6,r6,4                    ; r6 = add(values[j])
        slt  r9,r5,r1                   ; r9 = (j < nelem)
        bnez r9,loop2                   ; are still elements to be
                                           ; compared?

        nop
        nop
        addi r3,r3,1                     ; i = i + 1
        addi r2,r2,4                    ; r2 = add(values[i])
        j    loop1                      ; check for end of operations
        nop
        nop
end:    trap 0                          ; end of program

```

Taking into account now branch prediction, one has for *branch predictor* alternatives *none* or *static – predict always not taken*, which are the same for the DLX simulator, 620 instructions executed in 887 clock cycles yielding a speed up of

$$\text{speed up} = \frac{\text{clock cycles}_{\text{nonpipe}}}{\text{clock cycles}_{\text{pipe-nottaken}}} = \frac{2901}{887} = 3,27$$

On the other hand, for the *branch predictor* alternative *static – predict always taken*, there is no difference in this case.

4.3. Turn on the MIPS compatibility mode, in order to enable *pipeline interlocking*, and discard the remaining `nop` instructions. Check that your initial code now runs correctly. Go through the clock cycle diagram to understand why. Explain what has happened in your own words. What is the speed up now attained?

With the MIPS compatibility mode now turned on, not only all the `nop` instructions can be removed, but also one can take advantage of the branch delay slot feature.

```

.data
values: .word 1,2,3,4,5,6,7,8,9,10      ; values to be ordered
nelem:  .word 10                        ; array size
.text
.global main
main:   addi    r1,r0,nelem              ; r1 = add(nelem)
        lw      r1,0(r1)                ; r1 = val(nelem)
        addi    r2,r0,values            ; r2 = add(values[0])
        add     r3,r0,r0                ; r3 = i = 0 (counting variable)
        addi    r8,r1,-1                ; r8 = nelem - 1
loop1:  slt     r9,r3,r8                ; r9 = (i < nelem - 1)
        beqz    r9,end                 ; is the end of operations been
                                           ; reached?
        addi    r6,r2,4                 ; r6 = add(values[j])
        lw      r4,0(r2)                ; r4 = val(values[i])
        addi    r5,r3,1                 ; r5 = j = i+1 (counting variable)
        lw      r7,0(r6)                ; r7 = val(values[j])
loop2:  slt     r9,r4,r7                 ; r9 =
                                           ; (val(values[i]) < val(values[j]))
        beqz    r9,goon                 ; no element swap is required?
        add     r9,r4,r0                ; r9 = tmp = val(values[i])
        add     r4,r7,r0                ; val(values[i]) = val(values[j])
        add     r7,r9,r0                ; val(values[j]) = tmp
        sw      0(r2),r4                ; val(values[i]) = r4
        sw      0(r6),r7                ; val(values[j]) = r7
goon:   addi    r5,r5,1                 ; j = j + 1
        addi    r6,r6,4                 ; r6 = add(values[j])
        slt     r9,r5,r1                ; r9 = (j < nelem)
        bnez    r9,loop2                ; are still elements to be
                                           ; compared?
        lw      r7,0(r6)                ; r7 = val(values[j])
        addi    r3,r3,1                 ; i = i + 1
        j       loop1                  ; check for end of operations
        addi    r2,r2,4                 ; r2 = add(values[i])
end:    trap    0                       ; end of program

```

- 4.4. Consider the *branch predictor* alternatives: *none*, *static – predict always not taken* and *static – predict always taken*. Which one produces the most efficient run? Why? What is the speed up now attained?

In this case, both alternatives *static – predict always not taken* and *static – predict always taken* produce the same result: 620 instructions are executed in 689 clock cycles, yielding a speed up of

$$\text{speed up} = \frac{\text{clock cycles}_{\text{nonpipe}}}{\text{clock cycles}_{\text{pipe}}} = \frac{2901}{689} = 4,21 \text{ .}$$